

Part 4: Dimensionality Reduction

In this section, we'll perform dimensionality reduction with principal components analysis (PCA) and Linear Discriminant Analysis (LDA) on the 'Wine Rating and Price' dataset which can be found here. (Check out the classification section of this assignment to find out more about the wine data and what we will try to do with dimensionality reduction, as we'll make comparisons with the findings in that section.)

Wine Dataset

For PDF viewers: <https://www.kaggle.com/datasets/budnyak/wine-rating-and-price?select=Red.csv>

The wine data that was downloaded from kaggle was concatenated into a larger, total wine csv file. While there are distinct qualities in each type of wine that make them unique, we needed to concatenate them so that we have a large enough data set to perform analysis on. Additionally, all the wine categories held the same features, so a concatenation is not too messy to deal with.

```
wine <- read.csv("totalWine.csv", header = TRUE)
head(wine)
```

```
##              Name Country      Region      Winery
## 1          Pomerol 2011   France    Pomerol  Château La Providence
## 2             Lirac 2017   France      Lirac   Château Mont-Redon
## 3 Erta e China Rosso di Toscana 2015   Italy   Toscana       Renzo Masi
## 4             Bardolino 2019   Italy   Bardolino     Cavalchina
## 5      Ried Scheibner Pinot Noir 2016 Austria Carnuntum     Markowitsch
## 6  Gigondas (Nobles Terrasses) 2017   France  Gigondas     Vieux Clocher
##   Rating NumberOfRatings Price Year
## 1      4.2             100 95.00 2011
## 2      4.3             100 15.50 2017
## 3      3.9             100  7.45 2015
## 4      3.5             100  8.72 2019
## 5      3.9             100 29.15 2016
## 6      3.7             100 19.90 2017
```

```
tail(wine)
```

```
##              Name      Country      Region      Winery Rating
## 13829 Blanco (Verdejo) 2018      Spain      Rueda Marqués de Riscal    3.7
## 13830 Sauvignon Blanc 2019 New Zealand Marlborough     Oyster Bay    4.0
## 13831 Vinho Verde Sweet N.V. Portugal Vinho Verde     Casal Garcia    4.0
## 13832 Sauvignon Blanc 2018 New Zealand Marlborough     Kim Crawford    3.9
## 13833 Sauvignon Blanc 2019 New Zealand Marlborough     Hans Greyl    4.2
## 13834 Vinho Verde Branco N.V. Portugal Vinho Verde     Casal Garcia    3.7
##   NumberOfRatings Price Year
## 13829           4155  6.30 2018
## 13830           4423 10.66 2019
## 13831           4609  5.05 N.V.
## 13832           5105 14.90 2018
## 13833           5817  7.75 2019
## 13834          62980  4.35 N.V.
```

```
str(wine)
```

```
## 'data.frame':   13834 obs. of  8 variables:
##  $ Name      : chr  "Pomerol 2011" "Lirac 2017" "Erta e China Rosso di Toscana 2015" "Bardolino
##  $ Country   : chr  "France" "France" "Italy" "Italy" ...
##  $ Region    : chr  "Pomerol" "Lirac" "Toscana" "Bardolino" ...
##  $ Winery    : chr  "Château La Providence" "Château Mont-Redon" "Renzo Masi" "Cavalchina" ...
```

```
## $ Rating      : num  4.2 4.3 3.9 3.5 3.9 3.7 4 3.9 3.6 3.5 ...
## $ NumberOfRatings: int  100 100 100 100 100 100 100 100 100 100 ...
## $ Price       : num  95 15.5 7.45 8.72 29.15 ...
## $ Year        : chr   "2011" "2017" "2015" "2019" ...
```

Just for fun, let's find some max and min data from the super-dataset.

```
print(max(wine$Rating))
```

```
## [1] 4.9
```

```
print(min(wine$Rating))
```

```
## [1] 2.2
```

```
print(max(wine$Price))
```

```
## [1] 3410.79
```

```
print(min(wine$Price))
```

```
## [1] 3.15
```

We have some pretty nice wine here... and some pretty awful selections as well. Though these values don't tell us much. Let's take a closer look at some of these values.

```
print(subset(wine, Rating == max(Rating)))
```

```
##                               Name Country
## 10612 Montrachet Grand Cru Marquis de Laguiche 2017  France
##                               Region      Winery Rating NumberOfRatings  Price Year
## 10612 Montrachet Grand Cru Joseph Drouhin      4.9              34 681.37 2017
```

```
print(subset(wine, Rating == min(Rating)))
```

```
##                               Name Country      Region  Winery Rating
## 11948 Greca Terra Retsina N.V.  Greece Peloponnesos Tsantali    2.2
##                               NumberOfRatings Price Year
## 11948                               77  5.35 N.V.
```

```
print(subset(wine, Price == max(Price)))
```

```
##                               Name Country  Region Winery Rating NumberOfRatings  Price Year
## 2345 Pomerol 2012  France Pomerol Pétrus      4.7              204 3410.79 2012
```

```
print(subset(wine, Price == min(Price)))
```

```
##                               Name Country Region  Winery Rating
## 9164          Frizzantino Dolce N.V.  Italy Emilia Gualtieri    4.2
## 9325 Lambrusco dell'Emilia Dolce N.V.  Italy Emilia Gualtieri    3.8
##                               NumberOfRatings Price Year
## 9164                               43  3.15 N.V.
## 9325                               106  3.15 N.V.
```

It looks like this trivial data exploration had a purpose after all! Thanks to the Gualtieri (which looks to have two of the ultimate value choices of wine on the list) winery we find some null values in our year column. We'll need to find a way to deal with them to get a cleaner experiment done. For now, let's check how many of these values we have. We'll use the table function to see how many occurrences each value in the year column has.

```
nv_occur <- data.frame(table(wine$Year))
print(nv_occur)
```

```
##      Var1 Freq
## 1  1961     3
## 2  1988     1
## 3  1989     2
## 4  1990     2
## 5  1991     1
## 6  1992     3
## 7  1993     2
## 8  1995     4
## 9  1996     6
## 10 1997     7
## 11 1998     7
## 12 1999    19
## 13 2000    19
## 14 2001    12
## 15 2002    11
## 16 2003    15
## 17 2004    34
## 18 2005   160
## 19 2006    64
## 20 2007    60
## 21 2008   101
## 22 2009    99
## 23 2010   192
## 24 2011   312
## 25 2012   423
## 26 2013   624
## 27 2014   905
## 28 2015  1678
## 29 2016  2294
## 30 2017  2412
## 31 2018  2723
## 32 2019   893
## 33 2020     2
## 34 N.V.   744
```

Our year column seems to house a lot of “N.V.” values, 744 of them. This is concerning! Where are all these N.V.’s coming from!

Recall how the totalWine.csv file was constructed. We concatenated multiple types of wine (Red, Rose, Sparkling and White) that were all in separate .csv files.

```
spark <- read.csv("Sparkling.csv", header = TRUE)
nv_occur_spark <- data.frame(table(spark$Year))
print(nv_occur_spark)
```

```
##      Var1 Freq
## 1  1961     3
## 2  1996     1
## 3  1999     2
## 4  2002     4
## 5  2003     1
## 6  2004     5
## 7  2005     2
```

```
## 8 2006 15
## 9 2007 14
## 10 2008 18
## 11 2009 13
## 12 2010 13
## 13 2011 10
## 14 2012 25
## 15 2013 16
## 16 2014 19
## 17 2015 29
## 18 2016 22
## 19 2017 26
## 20 2018 28
## 21 2019 13
## 22 N.V. 728
```

As you can see, we have a bit of an issue with our “Sparkling” section of our wine. 728/744 null values come from that section. I cannot claim to be a domain expert on wine, but a quick google search tells me that wines made from a blend of other wines will exclude a vintage date, which can explain why a lot of these sparkling wines have no dates attached to them.

This concludes an already lengthy data exploration section.

We’ll copy the same conventions used in the classification section of the assignment, as we seek to make comparisons with accuracy after we complete dimensionality reduction.

Let’s divide into train and test.

```
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.2 --
## v ggplot2 3.3.6      v purrr 0.3.5
## v tibble 3.1.8       v dplyr 1.0.10
## v tidyr 1.2.1        v stringr 1.4.1
## v readr 2.1.3        v forcats 0.5.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()

library(dplyr)
library(ROCR)
library(mccr)
library(ISLR)
library(caret)

## Loading required package: lattice
##
## Attaching package: 'caret'
##
## The following object is masked from 'package:purrr':
##
##     lift

library(tree)
library(rpart)

red <- read.csv("Red.csv")
white <- read.csv("White.csv")
```

```

rose <- read.csv("Rose.csv")
sparkle <- read.csv("Sparkling.csv")
totalWine <- rbind(data = red, data = white, data = rose, data = sparkle)
names(totalWine)[1] <- "Name"

```

Now we need to clean our data.

```

totalWine <- subset(totalWine, select = -c(Name, Winery, Region))
totalWine <- subset(totalWine, totalWine$Year != "N.V.")
totalWine <- subset(totalWine, totalWine$Rating != 3)
totalWine <- subset(totalWine, totalWine$Year >= 2000)
totalWine$Rating[totalWine$Rating <= 3] <- 0
totalWine$Rating[totalWine$Rating > 3] <- 1
totalWine <- totalWine[,c(1,2,3,5,4)]

```

And finally our train/test split.

```

set.seed(512)
i <- sample(1 : nrow(totalWine), round(nrow(totalWine) * 0.8), replace = FALSE)
train <- totalWine[i,]
test <- totalWine[-i,]
str(train)

```

```

## 'data.frame':    10401 obs. of  5 variables:
## $ Country      : chr  "Austria" "Italy" "Italy" "France" ...
## $ Rating       : num  1 1 1 1 1 1 1 1 1 1 ...
## $ NumberOfRatings: num  30 28 375 40 290 42 791 110 193 282 ...
## $ Year         : chr  "2017" "2018" "2018" "2017" ...
## $ Price        : num  13.2 12.2 14.4 9.5 69.4 ...

```

PCA

Now that our data is properly split, we'll go ahead and perform PCA on our data sets. PCA is an unsupervised dimensionality reduction algorithm, as it pays no attention to class. It reduces the data and plots them on a new coordinate space while reducing the axes. These reduced axes are principal components.

```

pca_out <- preProcess(train[,1:5],method=c("center","scale","pca"))
pca_out

```

```

## Created from 10401 samples and 5 variables
##
## Pre-processing:
## - centered (3)
## - ignored (2)
## - principal component signal extraction (3)
## - scaled (3)
##
## PCA needed 3 components to capture 95 percent of the variance
summary(pca_out)

```

```

##           Length Class  Mode
## dim           2    -none- numeric
## bc            0    -none-  NULL
## yj            0    -none-  NULL
## et            0    -none-  NULL

```

```
## invHyperbolicSine 0      -none- NULL
## mean              3      -none- numeric
## std               3      -none- numeric
## ranges            0      -none- NULL
## rotation          9      -none- numeric
## method            4      -none- list
## thresh            1      -none- numeric
## pcaComp            0      -none- NULL
## numComp            1      -none- numeric
## ica                0      -none- NULL
## wildcards          2      -none- list
## k                  1      -none- numeric
## knnSummary         1      -none- function
## bagImp             0      -none- NULL
## median            0      -none- NULL
## data              0      -none- NULL
## rangeBounds        2      -none- numeric

prWine <- prcomp(~Rating + NumberOfRatings + Price, data = totalWine, scale = TRUE, center = TRUE)
prWine
```


```
## Standard deviations (1, ..., p=3):
## [1] 1.0172517 0.9945634 0.9879488
##
## Rotation (n x k) = (3 x 3):
##              PC1      PC2      PC3
## Rating      -0.4884176  0.8706445 -0.05853558
## NumberOfRatings -0.6114703 -0.3893384 -0.68885393
## Price        -0.6225370 -0.3006556  0.72253294
```

As we can see, PCA found three principal components that are responsible for 95% of the variance. I will elect to leave out PC3, as a vast majority of the variance is found in PC1 and PC2. (As the first PC captures the most variance, and subsequent PCs will represent reducing variances.) Additionally, you may have noticed that we only have one attribute that is continuous (Price), this is problematic as PCA would assume that it performs its methods on continuous variables when we've included two other discrete measures in Year and Rating. This may cause some trouble with our predictions.


```
train_pc <- predict(pca_out, train[, 1:5])
summary(train_pc)
```

```
##      Country      Year      PC1      PC2
## Length:10401    Length:10401    Min.   :-16.54772    Min.   :-20.7843
## Class :character Class :character 1st Qu.: -0.06901    1st Qu.:  0.0163
## Mode  :character Mode  :character Median :  0.24843    Median :  0.1735
##                                     Mean  :  0.00000    Mean   :  0.0000
##                                     3rd Qu.:  0.39359    3rd Qu.:  0.2441
##                                     Max.   : 10.95378    Max.   :  0.3122
##      PC3
## Min.   :-18.48415
## 1st Qu.: -0.15550
## Median : -0.01497
## Mean   :  0.00000
## 3rd Qu.:  0.13977
## Max.   : 18.91313
```

```
test_pc <- predict(pca_out, test[,])
plot(test_pc$PC1, test_pc$PC2)
```



```
plot(test_pc$PC1, test_pc$PC2, pch=c(23,21,22), bg=c("red","green","blue"))
```



Now

that we have done our PCA, let's try linear regression.

```
lm1 <- lm(train_pc$PC1 ~ ., data = train_pc)
summary(lm1)
```

```
##
## Call:
```

```
## lm(formula = train_pc$PC1 ~ ., data = train_pc)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -16.8077  -0.0996   0.1177   0.3170  10.0909
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -1.947571    0.247710  -7.862 4.15e-15 ***
## CountryAustralia    0.151215    0.086921   1.740 0.081944 .
## CountryAustria     0.549253    0.079132   6.941 4.13e-12 ***
## CountryBrazil      1.093705    0.173429   6.306 2.97e-10 ***
## CountryBulgaria    0.496977    0.677377   0.734 0.463161
## CountryCanada      0.707665    0.955950   0.740 0.459151
## CountryChile       0.222555    0.079948   2.784 0.005383 **
## CountryChina       3.437206    0.557904   6.161 7.50e-10 ***
## CountryCroatia     0.596618    0.481002   1.240 0.214868
## CountryCzech Republic 0.375018    0.677401   0.554 0.579856
## CountryFrance      0.241890    0.065233   3.708 0.000210 ***
## CountryGeorgia     0.615464    0.308142   1.997 0.045814 *
## CountryGermany     0.491436    0.069275   7.094 1.39e-12 ***
## CountryGreece      0.607362    0.227487   2.670 0.007600 **
## CountryHungary     0.682163    0.254011   2.686 0.007252 **
## CountryIsrael      0.422327    0.227414   1.857 0.063327 .
## CountryItaly       0.313732    0.064231   4.884 1.05e-06 ***
## CountryLebanon     -0.519928    0.300031  -1.733 0.083141 .
## CountryLuxembourg   0.624488    0.480926   1.299 0.194140
## CountryMexico      0.232642    0.956626   0.243 0.807863
## CountryMoldova     -0.062071    0.307907  -0.202 0.840241
## CountryNew Zealand  0.225620    0.104683   2.155 0.031164 *
## CountryPortugal    0.413266    0.086162   4.796 1.64e-06 ***
## CountryRomania     0.675437    0.184928   3.652 0.000261 ***
## CountrySlovakia    0.639328    0.677386   0.944 0.345286
## CountrySlovenia    0.631912    0.282277   2.239 0.025201 *
## CountrySouth Africa 0.468376    0.071791   6.524 7.16e-11 ***
## CountrySpain       0.327171    0.067503   4.847 1.27e-06 ***
## CountrySwitzerland 0.356522    0.227498   1.567 0.117111
## CountryTurkey      0.568930    0.324048   1.756 0.079170 .
## CountryUnited Kingdom 0.245637    0.554625   0.443 0.657856
## CountryUnited States -0.167760    0.077213  -2.173 0.029826 *
## CountryUruguay     -0.532162    0.554220  -0.960 0.336978
## Year2001           1.170203    0.398062   2.940 0.003292 **
## Year2002          -0.265421    0.432899  -0.613 0.539807
## Year2003           0.553260    0.364327   1.519 0.128899
## Year2004           0.618371    0.299050   2.068 0.038685 *
## Year2005           0.664562    0.254822   2.608 0.009122 **
## Year2006           0.192919    0.276723   0.697 0.485719
## Year2007           0.884881    0.276914   3.196 0.001400 **
## Year2008           0.612363    0.262004   2.337 0.019446 *
## Year2009           0.644891    0.263526   2.447 0.014415 *
## Year2010           0.664953    0.251239   2.647 0.008141 **
## Year2011           1.276985    0.246814   5.174 2.34e-07 ***
## Year2012           1.321315    0.245142   5.390 7.20e-08 ***
## Year2013           1.383877    0.243527   5.683 1.36e-08 ***
```

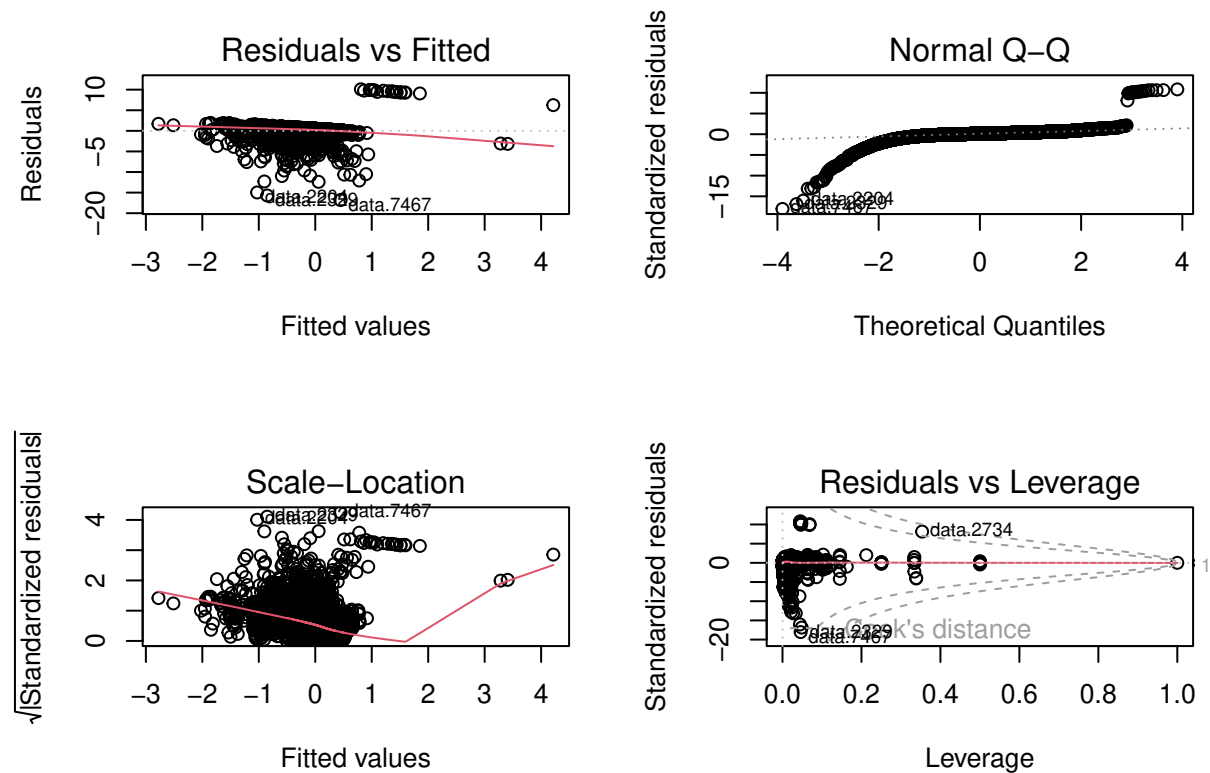


```
## Year2014          1.505250    0.242455    6.208 5.56e-10 ***
## Year2015          1.542855    0.241325    6.393 1.69e-10 ***
## Year2016          1.650301    0.240959    6.849 7.87e-12 ***
## Year2017          1.782088    0.241103    7.391 1.57e-13 ***
## Year2018          1.936895    0.241028    8.036 1.03e-15 ***
## Year2019          1.994149    0.242817    8.213 2.42e-16 ***
## Year2020          1.966979    0.716857    2.744 0.006082 **
## PC2               -0.044483    0.009606   -4.631 3.68e-06 ***
## PC3                0.064372    0.009998    6.439 1.26e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.9537 on 10346 degrees of freedom
## Multiple R-squared:  0.1316, Adjusted R-squared:  0.127
## F-statistic: 29.03 on 54 and 10346 DF,  p-value: < 2.2e-16
```

Now let's plot the residuals.

```
par(mfrow = c(2,2))
plot(lm1)
```

```
## Warning: not plotting observations with leverage one:
##      2844
## Warning in sqrt(crit * p * (1 - hh)/hh): NaNs produced
## Warning in sqrt(crit * p * (1 - hh)/hh): NaNs produced
```



```
pred1 <- predict(lm1, newdata = test_pc)
summary(pred1)
```

```
##      Min.      1st Qu.      Median      Mean      3rd Qu.      Max.
```

```
## -1.938770 -0.156253 0.038845 -0.006681 0.262473 2.811055
```

```
print(cor(pred1, test_pc$PC1))
```

```
## [1] 0.1674418
```

```
print(mean(pred1 - test_pc$PC1) ^ 2)
```

```
## [1] 7.671703e-06
```

```
rmse1 <- sqrt(mean(pred1 - test_pc$PC1) ^ 2)
```

```
rmse1
```

```
## [1] 0.002769784
```

We can see that our correlation values fell a bit, signaling a slight reduction in accuracy. Which makes sense as we had reduced our data.

LDA

Linear discriminant analysis WILL consider class, unlike its cousin PCA. Thus, LDA can be superior when we have a known class to fit data into. LDA finds a linear combination of predictors that maximizes separation between classes and minimizes the standard deviation within a class.

```
library(MASS)
```

```
##
```

```
## Attaching package: 'MASS'
```

```
## The following object is masked from 'package:dplyr':
```

```
##
```

```
##      select
```

```
lda1 <- lda(Price~., data = train)
```

```
lda1$means
```

| ## | CountryAustralia | CountryAustria | CountryBrazil | CountryBulgaria |
|---------|------------------|----------------|---------------|-----------------|
| ## 3.55 | 0.00000000 | 0.00000000 | 0.00000000 | 0.00000000 |
| ## 3.7 | 0.00000000 | 0.00000000 | 0.00000000 | 0.00000000 |
| ## 3.74 | 0.00000000 | 0.00000000 | 0.00000000 | 0.00000000 |
| ## 3.75 | 0.00000000 | 0.00000000 | 0.00000000 | 0.00000000 |
| ## 3.79 | 0.00000000 | 0.00000000 | 0.00000000 | 0.00000000 |
| ## 3.95 | 0.00000000 | 0.00000000 | 0.00000000 | 0.00000000 |
| ## 3.99 | 0.00000000 | 0.00000000 | 0.00000000 | 0.00000000 |
| ## 4 | 0.00000000 | 0.00000000 | 0.00000000 | 0.00000000 |
| ## 4.15 | 0.00000000 | 0.00000000 | 0.00000000 | 0.00000000 |
| ## 4.16 | 0.00000000 | 0.00000000 | 0.00000000 | 0.00000000 |
| ## 4.25 | 0.00000000 | 0.00000000 | 0.00000000 | 0.00000000 |
| ## 4.28 | 0.00000000 | 0.00000000 | 0.00000000 | 0.00000000 |
| ## 4.29 | 0.00000000 | 0.00000000 | 0.00000000 | 0.00000000 |
| ## 4.3 | 0.00000000 | 0.00000000 | 0.00000000 | 0.00000000 |
| ## 4.31 | 0.00000000 | 0.00000000 | 0.00000000 | 0.00000000 |
| ## 4.38 | 0.00000000 | 0.00000000 | 0.00000000 | 0.00000000 |
| ## 4.4 | 0.00000000 | 0.00000000 | 0.00000000 | 0.00000000 |
| ## 4.44 | 0.00000000 | 0.00000000 | 0.00000000 | 0.00000000 |
| ## 4.45 | 0.00000000 | 0.00000000 | 0.00000000 | 0.00000000 |
| ## 4.5 | 0.00000000 | 0.00000000 | 0.00000000 | 0.00000000 |
| ## 4.55 | 0.00000000 | 0.00000000 | 0.00000000 | 0.00000000 |

[illegible]

[illegible]

| | | | | |
|--------------------------|---------------|---------------|---------------|---------------|
| ## CountryLebanon | -6.787403e-02 | 4.172415e-02 | -1.868567e-02 | 3.909762e-02 |
| ## CountryLuxembourg | -2.704892e-02 | 8.856544e-02 | 5.681759e-03 | -5.250638e-02 |
| ## CountryMexico | -4.400869e-02 | 2.745469e-01 | 4.512722e-02 | -1.461833e-01 |
| ## CountryMoldova | -8.066841e-04 | 3.709996e-02 | -2.759620e-03 | -4.596180e-02 |
| ## CountryNew Zealand | 1.511986e-01 | -1.240325e-01 | 5.884640e-02 | -1.357162e-01 |
| ## CountryPortugal | -3.593999e-01 | 9.975121e-02 | -2.041191e-02 | -9.821181e-03 |
| ## CountryRomania | 5.381921e-01 | 1.496702e+01 | -2.093963e+00 | -5.114234e+00 |
| ## CountrySlovakia | 1.017742e+01 | 2.115199e+01 | 1.649101e+01 | 5.629167e+01 |
| ## CountrySlovenia | -1.545114e+01 | 5.468643e-01 | 1.279180e-01 | 1.307837e+00 |
| ## CountrySouth Africa | -1.687347e-02 | -1.377714e-01 | 6.983389e-02 | -1.157027e-01 |
| ## CountrySpain | 1.381214e-02 | 4.197080e-02 | 1.440170e-02 | -1.981445e-02 |
| ## CountrySwitzerland | -4.758196e-02 | -6.645005e-02 | 1.635612e-02 | 2.324576e-02 |
| ## CountryTurkey | -3.918113e+00 | -5.327550e+00 | 6.140311e-01 | 2.011485e+00 |
| ## CountryUnited Kingdom | -1.613954e-01 | 1.692158e-02 | -2.639647e-02 | 4.376702e-02 |
| ## CountryUnited States | 1.009736e-02 | -1.270576e-02 | 3.299893e-02 | -2.269264e-02 |
| ## CountryUruguay | 2.058015e+00 | 5.830943e-01 | -1.785684e-02 | 1.687805e-01 |
| ## Rating | -1.585820e-01 | 2.854437e-01 | 2.956890e-02 | -1.394051e-01 |
| ## NumberOfRatings | 6.344999e-06 | -7.961667e-06 | 6.032259e-06 | -3.271830e-06 |
| ## Year2001 | -3.816442e-02 | 8.939072e-03 | 2.504713e-03 | 1.387763e-02 |
| ## Year2002 | -1.446647e-02 | 1.310904e-03 | -6.480272e-03 | 6.999415e-03 |
| ## Year2003 | 4.525561e-02 | 1.073489e-01 | -3.577545e-03 | -3.193922e-02 |
| ## Year2004 | -1.790243e-02 | -5.212768e-03 | 1.047114e-02 | -2.709072e-03 |
| ## Year2005 | -2.473786e-02 | 9.863940e-03 | -1.232727e-02 | 1.420073e-02 |
| ## Year2006 | -5.968087e-03 | 1.955486e-03 | 2.529429e-03 | 3.054409e-03 |
| ## Year2007 | 1.557529e-03 | 1.905899e-02 | -5.753306e-03 | 4.333951e-03 |
| ## Year2008 | 3.543613e-02 | 1.442041e-02 | -2.189123e-02 | 5.185512e-03 |
| ## Year2009 | 5.612026e-02 | -8.708113e-03 | -1.508398e-02 | 1.529873e-02 |
| ## Year2010 | -2.867112e-02 | 1.458319e-02 | -1.264813e-02 | -1.797924e-02 |
| ## Year2011 | 3.562875e-02 | 1.598819e-02 | 4.142401e-03 | -4.502524e-02 |
| ## Year2012 | 3.155637e-02 | 3.473633e-02 | -2.707898e-03 | -9.261088e-03 |
| ## Year2013 | 6.469553e-02 | 2.945535e-03 | -7.661883e-03 | -4.013230e-02 |
| ## Year2014 | 5.722172e-02 | 2.471491e-02 | 7.339117e-03 | -2.829255e-02 |
| ## Year2015 | -2.736074e-02 | 2.449592e-02 | 4.385395e-02 | -4.777540e-02 |
| ## Year2016 | 2.099765e-03 | -2.377244e-02 | -2.187815e-03 | -5.874881e-03 |
| ## Year2017 | 3.162700e-02 | -8.278620e-02 | 1.043988e-02 | 1.028478e-02 |
| ## Year2018 | -5.897151e-03 | -6.109832e-02 | 2.134442e-02 | 3.431160e-02 |
| ## Year2019 | 4.314965e-02 | -4.247977e-02 | -1.300157e-02 | 1.096318e-02 |
| ## Year2020 | -1.827492e-02 | 2.319103e+00 | -1.000807e+00 | -1.224477e+00 |
| ## | LD53 | LD54 | | |
| ## CountryAustralia | -2.639318e-02 | -8.469747e-04 | | |
| ## CountryAustria | -4.254561e-02 | 4.696618e-02 | | |
| ## CountryBrazil | 8.474102e-03 | 1.140675e-01 | | |
| ## CountryBulgaria | -1.779625e-02 | 3.236773e-03 | | |
| ## CountryCanada | -6.280130e-02 | -5.827562e-01 | | |
| ## CountryChile | -1.806197e-03 | 4.084109e-02 | | |
| ## CountryChina | -2.706210e-03 | 5.098557e-02 | | |
| ## CountryCroatia | 3.452416e+01 | 4.640353e+00 | | |
| ## CountryCzech Republic | -3.816555e+01 | -5.049928e+00 | | |
| ## CountryFrance | -4.629070e-02 | 1.450437e-02 | | |
| ## CountryGeorgia | -2.627039e-03 | 5.555491e-02 | | |
| ## CountryGermany | -3.170398e-02 | 1.660121e-02 | | |
| ## CountryGreece | 1.222406e+00 | 3.616305e-01 | | |
| ## CountryHungary | -9.920278e-01 | 3.979240e-02 | | |
| ## CountryIsrael | 2.337067e-02 | 9.713556e-02 | | |

```
## CountryItaly -2.011163e-02 8.057530e-03
## CountryLebanon -1.644750e-02 -2.232072e-03
## CountryLuxembourg 3.403336e-03 2.963174e-02
## CountryMexico 3.008343e-03 1.771995e-01
## CountryMoldova -2.229544e-02 2.285581e-02
## CountryNew Zealand -5.762401e-02 6.604469e-02
## CountryPortugal -6.333612e-03 1.715160e-02
## CountryRomania -9.793819e-01 -7.624314e-01
## CountrySlovakia 3.089524e+00 1.158605e+00
## CountrySlovenia 3.951880e-01 2.300480e-01
## CountrySouth Africa -7.885672e-03 3.057292e-02
## CountrySpain -1.490683e-02 3.096824e-02
## CountrySwitzerland -4.966276e-03 5.185219e-02
## CountryTurkey 3.415580e-01 4.412937e-01
## CountryUnited Kingdom 4.329752e-02 1.543064e-02
## CountryUnited States -2.069601e-03 1.681490e-02
## CountryUruguay -3.384096e+00 -3.025339e-01
## Rating -1.144184e-01 5.404661e-01
## NumberOfRatings -5.374291e-06 8.479286e-06
## Year2001 7.256051e-03 -2.248157e-03
## Year2002 3.611548e-03 -3.296748e-03
## Year2003 4.948091e-03 -4.626497e-03
## Year2004 2.712387e-03 4.321348e-03
## Year2005 5.149756e-03 -1.523470e-03
## Year2006 2.077651e-03 9.355211e-04
## Year2007 4.126353e-03 1.113096e-02
## Year2008 4.609467e-02 3.971207e-03
## Year2009 -7.496749e-02 -6.902149e-03
## Year2010 -1.210951e-02 3.869984e-03
## Year2011 -1.613261e-02 -1.684513e-02
## Year2012 1.059868e-03 6.754828e-03
## Year2013 -2.417498e-02 6.475244e-03
## Year2014 -1.843354e-03 1.066015e-02
## Year2015 -7.477242e-03 -4.137931e-04
## Year2016 -2.124870e-02 3.678406e-03
## Year2017 -5.526997e-02 5.553615e-03
## Year2018 3.306930e-03 -1.827223e-02
## Year2019 -1.208776e-02 -1.504149e-01
## Year2020 -8.626739e+00 6.230472e+01
```

Now let's test our LDA1 model.

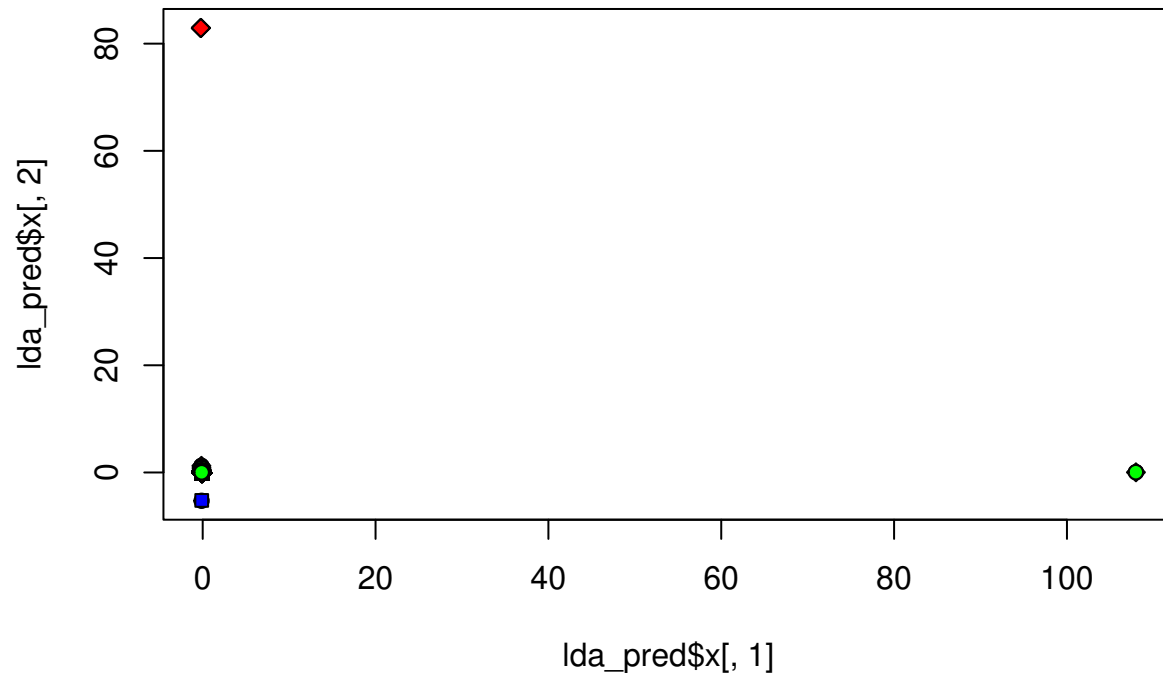
```
lda_pred <- predict(lda1, newdata = test, type = "class")
options(max.print = 100) #the output here was excessive, so I had to cut it down slightly.
lda_pred$class
```

```
## [1] 129.01 32.14 15.5 30.74 15.5 140.64 43.16 32.75 8.9 54.5
## [11] 5.36 32.75 15.5 75.07 7.42 37.73 161.1 120.01 156.17 36.89
## [21] 9.34 43.16 58.8 7.37 36.83 7.84 100 45.77 5.36 130.97
## [31] 36.83 15.5 5.31 129.01 26.06 18.4 484.03 19.24 7.65 63.26
## [41] 15.5 43.16 32.75 7.42 5.31 25.84 328.95 5.99 58.8 7.42
## [51] 129.01 62.39 62.39 328.95 29.1 20.19 36.63 139.95 129.01 11.56
## [61] 130.97 28.6 43.16 39.89 125.65 28.41 7.42 12.98 208.44 5.36
## [71] 43.16 5.36 7.84 36.63 22.85 20.19 7.42 54.59 32.75 5.99
## [81] 29.1 9.34 7.84 23.38 18.4 3.74 137.58 137.58 29.12 76.76
```

```
## [91] 63.66 130.97 15.5 721.34 19.05 8.9 484.03 484.03 17.42 129.01
## [ reached getOption("max.print") -- omitted 2500 entries ]
## 2616 Levels: 3.55 3.7 3.74 3.75 3.79 3.95 3.99 4 4.15 4.16 4.25 4.28 4.29 ... 1599.95
mean(lda_pred$class == test$Price)
```

```
## [1] 0.005
```

```
plot(lda_pred$x[,1], lda_pred$x[,2], pch=c(23,21,22), bg=c("red","green","blue"))
```



With a mean of 0.005, we have an absolutely diabolical accuracy score.