

In []:

In []:

Image Classification

"Atmin Sheth" and Abigail Solomon

1. Import Libraries

```
In [1]: import pandas as pd
import numpy as np
import cv2
import seaborn as sns
import matplotlib.pyplot as plt
import keras
import os
import tensorflow as tf
from keras.preprocessing.image import ImageDataGenerator
from keras.models import Sequential
from keras import layers
from keras.layers import Conv2D, MaxPooling2D
from keras.layers import Activation, Dropout, Flatten, Dense
from keras.callbacks import EarlyStopping
from tensorflow.keras.optimizers import RMSprop
from sklearn.metrics import accuracy_score
```

2. Import dataset

```
In [2]: # Use google.colab to directly import dataset from Kaggle to colab with out do
wnloading to pc.
from google.colab import drive
drive.mount('/content/gdrive')
```

Mounted at /content/gdrive

```
In [3]: # Directory to dataset in drive
os.environ['KAGGLE_CONFIG_DIR'] = "/content/gdrive/My Drive/Kaggle"
```

```
In [4]: # Change directory
%cd /content/gdrive/My Drive/Kaggle

/content/gdrive/My Drive/Kaggle
```

```
In [5]: # Kaggle Dataset API
!kaggle datasets download -d hasibalmuzdadid/shoe-vs-sandal-vs-boot-dataset-15k-images

shoe-vs-sandal-vs-boot-dataset-15k-images.zip: Skipping, found more recently modified local copy (use --force to force download)
```

```
In [ ]: #Unzip folder and remove the zipped one
!unzip \*.zip && rm *.zip
```

3. Data Exploration

```
In [6]: from PIL import Image
rootDIR = "/content/gdrive/MyDrive/Kaggle/Shoe vs Sandal vs Boot Dataset"
ImageDIR = os.path.join(rootDIR)
ImageNames = os.listdir(ImageDIR)
TargetImageCount = len(ImageNames)
print('There are', TargetImageCount, 'image targets, and their names are the following:\n', ImageNames)
```

There are 3 image targets, and their names are the following:
['Boot', 'Sandal', 'Shoe']

```
In [7]: BootDIR = "/content/gdrive/MyDrive/Kaggle/Shoe vs Sandal vs Boot Dataset/Boot"
ImageDIR = os.path.join(BootDIR)
ImageFiles = os.listdir(ImageDIR)
ImageCount = len(ImageFiles)
print('There are', ImageCount, 'Boot images ')
```

There are 5000 Boot images

```
In [ ]: SandalDIR = "/content/gdrive/MyDrive/Kaggle/Shoe vs Sandal vs Boot Dataset/Sandal"
ImageDIR = os.path.join(SandalDIR)
ImageFiles = os.listdir(ImageDIR)
ImageCount = len(ImageFiles)
print('There are', ImageCount, 'Sandal images ')
```

There are 5000 Sandal images

```
In [8]: ShoeDIR = "/content/gdrive/MyDrive/Kaggle/Shoe vs Sandal vs Boot Dataset/Shoe"
ImageDIR = os.path.join(ShoeDIR)
ImageFiles = os.listdir(ImageDIR)
ImageCount = len(ImageFiles)
print('There are', ImageCount, 'Shoe images ')
```

There are 5000 Shoe images

```
In [9]: import tensorflow as tf
        from tensorflow import keras

        image_size = 256
        batch_size = 10
        epochs = 10

        target_size = (image_size, image_size)
        input_shape = (image_size, image_size, 3)
```

```
In [10]: # Loading sample images to visualize data

img = cv2.imread('/content/gdrive/MyDrive/Kaggle/Shoe vs Sandal vs Boot Datas
t/Boot/boot (100).jpg')
# converting to the RGB ordering that matplotlib wants
img_show1 = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

fig, ax = plt.subplots(1,2, figsize=(12,6))
ax[0].imshow(img) # RGB order from imread()
ax[0].axis('off')
ax[1].imshow(img_show1) # expected color to the image
ax[1].axis('off')
plt.show()

img2 = cv2.imread('/content/gdrive/MyDrive/Kaggle/Shoe vs Sandal vs Boot Datas
et/Sandal/Sandal (1003).jpg')
# converting to the RGB ordering that matplotlib wants
img_show2 = cv2.cvtColor(img2, cv2.COLOR_BGR2RGB)

fig, ax = plt.subplots(1,2, figsize=(12,6))
ax[0].imshow(img2) # RGB order from imread()
ax[0].axis('off')
ax[1].imshow(img_show2) # expected color to the image
ax[1].axis('off')
plt.show()

img3 = cv2.imread('/content/gdrive/MyDrive/Kaggle/Shoe vs Sandal vs Boot Datas
et/Shoe/Shoe (1004).jpg')
# converting to the RGB ordering that matplotlib wants
img_show3 = cv2.cvtColor(img3, cv2.COLOR_BGR2RGB)

fig, ax = plt.subplots(1,2, figsize=(12,6))
ax[0].imshow(img3) # RGB order from imread()
ax[0].axis('off')
ax[1].imshow(img_show3) # expected color to the image
ax[1].axis('off')
plt.show()
```



4. Split data into train and test

```
In [11]: # By setting validation_split to 0.2, 20% of the original data set will be generated for test and the remaining for train
# Training in TensorFlow using ImageDataGenerator, will automatically label images based on their parent directory.

datagen = ImageDataGenerator(rescale=1./255,rotation_range=45, horizontal_flip=True,
                             vertical_flip=True, fill_mode='reflect',validation_split=.20)

train = datagen.flow_from_directory(rootDIR, batch_size = 10,target_size=(256,256),
                                   classes = [ 'Boot','Sandal', 'Shoe'],
                                   class_mode='categorical', subset="training")

test = datagen.flow_from_directory(rootDIR, batch_size = 10,target_size=(256,256),
                                   classes = [ 'Boot','Sandal', 'Shoe'],
                                   class_mode='categorical', subset="validation")

Found 12000 images belonging to 3 classes.
Found 3000 images belonging to 3 classes.
```


```
In [12]: # The Classes are represented in numbers, since computers perform better with numbers, and also to avoid language bias
print (train.class_indices)

{'Boot': 0, 'Sandal': 1, 'Shoe': 2}
```

5. Bar graph

The dataset has 15,000 observations, and the tree target classes has 5000 observations each, the bar graph is displaying the distribution of the target classes visually.

```
In [13]: import matplotlib.pyplot as plt
fig = plt.figure()
ax = fig.add_axes([0,0,1,1])
ShoeImages = ['Boot', 'Sandal', 'Shoe']
ImageCount = [5000,5000,5000]
ax.bar(ShoeImages,ImageCount)
plt.show()
```



6. Description of the dataset

The dataset has a total of 15,000 images, labeled into three subclasses, named Boot, Sandal, and Shoe. Just like a human being, able to recognize anything after observing or using something many times, a machine can be trained to detect an image by feeding a large number of data to observe. Now, we will train a neural network with 80% of the data to learn whether the given image is Boot, Sandal, or Shoe, by first detecting the patterns in raw pixels to classify images, and then detect features using convolutions trained to spot particular features that make up a shoe. Then the trained model would be able to predict on the rest 20% of the data, it will classify the new data into Boot, Sandal or Shoe class in the same pattern.

```
In [14]: #Avoid overfit by making the process to terminate early
callback = keras.callbacks.EarlyStopping(monitor='accuracy', patience=10)
```

7. Create sequential model

```
In [15]: # Image recognition by processing pixel data
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(256, 256)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    # Three output neurons for classes
    tf.keras.layers.Dense(3, activation='softmax')
])
```

```
In [16]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 65536)	0
dense (Dense)	(None, 128)	8388736
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 128)	16512
dropout_1 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 3)	387
Total params: 8,405,635		
Trainable params: 8,405,635		
Non-trainable params: 0		

7. Evaluate on the test data

Compile and fit.

```
In [ ]: model.compile(loss='categorical_crossentropy',
                      optimizer='adam',
                      metrics=['accuracy'])
```

```
In [22]: total_sample = train.n
         n_epochs = 10
```



```
In [23]: history = model.fit(  
    x= test,  
    batch_size=10,  
    steps_per_epoch = int(total_sample/batch_size),  
    epochs=n_epochs,  
    verbose=1  
)
```

Epoch 1/10

WARNING:tensorflow:Model was constructed with shape (None, 256, 256) for input KerasTensor(type_spec=TensorSpec(shape=(None, 256, 256), dtype=tf.float32, name='flatten_input'), name='flatten_input', description="created by layer 'flatten_input'"), but it was called on an input with incompatible shape (None, None, None).

WARNING:tensorflow:Model was constructed with shape (None, 256, 256) for input KerasTensor(type_spec=TensorSpec(shape=(None, 256, 256), dtype=tf.float32, name='flatten_input'), name='flatten_input', description="created by layer 'flatten_input'"), but it was called on an input with incompatible shape (None, None, None).

```

-----
InvalidArgumentError                                Traceback (most recent call last)
<ipython-input-23-751ede4a5913> in <module>
----> 1 history = model.fit(
      2     x= test,
      3     batch_size=10,
      4     steps_per_epoch = int(total_sample/batch_size),
      5     epochs=n_epochs,

/usr/local/lib/python3.8/dist-packages/keras/utils/traceback_utils.py in error
r_handler(*args, **kwargs)
      65     except Exception as e: # pylint: disable=broad-except
      66         filtered_tb = _process_traceback_frames(e.__traceback__)
----> 67         raise e.with_traceback(filtered_tb) from None
      68     finally:
      69         del filtered_tb

/usr/local/lib/python3.8/dist-packages/tensorflow/python/eager/execute.py in
quick_execute(op_name, num_outputs, inputs, attrs, ctx, name)
      52     try:
      53         ctx.ensure_initialized()
----> 54         tensors = pywrap_tfe.TFE_Py_Execute(ctx._handle, device_name, op_
name,
      55                                         inputs, attrs, num_outputs)
      56     except core._NotOkStatusException as e:

InvalidArgumentError: Graph execution error:

Detected at node 'sequential/dense/Relu' defined at (most recent call last):
  File "/usr/lib/python3.8/runpy.py", line 194, in _run_module_as_main
    return _run_code(code, main_globals, None,
  File "/usr/lib/python3.8/runpy.py", line 87, in _run_code
    exec(code, run_globals)
  File "/usr/local/lib/python3.8/dist-packages/ipykernel_launcher.py", line
16, in <module>
    app.launch_new_instance()
  File "/usr/local/lib/python3.8/dist-packages/traitlets/config/applicatio
n.py", line 985, in launch_instance
    app.start()
  File "/usr/local/lib/python3.8/dist-packages/ipykernel/kernelapp.py", lin
e 612, in start
    self.io_loop.start()
  File "/usr/local/lib/python3.8/dist-packages/tornado/platform/asyncio.p
y", line 149, in start
    self.asyncio_loop.run_forever()
  File "/usr/lib/python3.8/asyncio/base_events.py", line 570, in run_foreve
r
    self._run_once()
  File "/usr/lib/python3.8/asyncio/base_events.py", line 1859, in _run_once
    handle._run()
  File "/usr/lib/python3.8/asyncio/events.py", line 81, in _run
    self._context.run(self._callback, *self._args)
  File "/usr/local/lib/python3.8/dist-packages/tornado/ioloop.py", line 69
0, in <lambda>
    lambda f: self._run_callback(functools.partial(callback, future))
  File "/usr/local/lib/python3.8/dist-packages/tornado/ioloop.py", line 74
3, in _run_callback

```

```

        ret = callback()
    File "/usr/local/lib/python3.8/dist-packages/tornado/gen.py", line 787, i
n inner
        self.run()
    File "/usr/local/lib/python3.8/dist-packages/tornado/gen.py", line 748, i
n run
        yielded = self.gen.send(value)
    File "/usr/local/lib/python3.8/dist-packages/ipykernel/kernelbase.py", li
ne 365, in process_one
        yield gen.maybe_future(dispatch(*args))
    File "/usr/local/lib/python3.8/dist-packages/tornado/gen.py", line 209, i
n wrapper
        yielded = next(result)
    File "/usr/local/lib/python3.8/dist-packages/ipykernel/kernelbase.py", li
ne 268, in dispatch_shell
        yield gen.maybe_future(handler(stream, idents, msg))
    File "/usr/local/lib/python3.8/dist-packages/tornado/gen.py", line 209, i
n wrapper
        yielded = next(result)
    File "/usr/local/lib/python3.8/dist-packages/ipykernel/kernelbase.py", li
ne 543, in execute_request
        self.do_execute(
    File "/usr/local/lib/python3.8/dist-packages/tornado/gen.py", line 209, i
n wrapper
        yielded = next(result)
    File "/usr/local/lib/python3.8/dist-packages/ipykernel/ipkernel.py", line
306, in do_execute
        res = shell.run_cell(code, store_history=store_history, silent=silent)
    File "/usr/local/lib/python3.8/dist-packages/ipykernel/zmqshell.py", line
536, in run_cell
        return super(ZMQInteractiveShell, self).run_cell(*args, **kwargs)
    File "/usr/local/lib/python3.8/dist-packages/IPython/core/interactiveshel
l.py", line 2854, in run_cell
        result = self._run_cell(
    File "/usr/local/lib/python3.8/dist-packages/IPython/core/interactiveshel
l.py", line 2881, in _run_cell
        return runner(coro)
    File "/usr/local/lib/python3.8/dist-packages/IPython/core/async_helpers.p
y", line 68, in _pseudo_sync_runner
        coro.send(None)
    File "/usr/local/lib/python3.8/dist-packages/IPython/core/interactiveshel
l.py", line 3057, in run_cell_async
        has_raised = await self.run_ast_nodes(code_ast.body, cell_name,
    File "/usr/local/lib/python3.8/dist-packages/IPython/core/interactiveshel
l.py", line 3249, in run_ast_nodes
        if (await self.run_code(code, result, async_=asy)):
    File "/usr/local/lib/python3.8/dist-packages/IPython/core/interactiveshel
l.py", line 3326, in run_code
        exec(code_obj, self.user_global_ns, self.user_ns)
    File "<ipython-input-23-751ede4a5913>", line 1, in <module>
        history = model.fit(
    File "/usr/local/lib/python3.8/dist-packages/keras/utils/traceback_utils.
py", line 64, in error_handler
        return fn(*args, **kwargs)
    File "/usr/local/lib/python3.8/dist-packages/keras/engine/training.py", l
ine 1409, in fit
        tmp_logs = self.train_function(iterator)

```

```

File "/usr/local/lib/python3.8/dist-packages/keras/engine/training.py", 1
ine 1051, in train_function
    return step_function(self, iterator)
File "/usr/local/lib/python3.8/dist-packages/keras/engine/training.py", 1
ine 1040, in step_function
    outputs = model.distribute_strategy.run(run_step, args=(data,))
File "/usr/local/lib/python3.8/dist-packages/keras/engine/training.py", 1
ine 1030, in run_step
    outputs = model.train_step(data)
File "/usr/local/lib/python3.8/dist-packages/keras/engine/training.py", 1
ine 889, in train_step
    y_pred = self(x, training=True)
File "/usr/local/lib/python3.8/dist-packages/keras/utils/traceback_utils.
py", line 64, in error_handler
    return fn(*args, **kwargs)
File "/usr/local/lib/python3.8/dist-packages/keras/engine/training.py", 1
ine 490, in __call__
    return super().__call__(*args, **kwargs)
File "/usr/local/lib/python3.8/dist-packages/keras/utils/traceback_utils.
py", line 64, in error_handler
    return fn(*args, **kwargs)
File "/usr/local/lib/python3.8/dist-packages/keras/engine/base_layer.py",
line 1014, in __call__
    outputs = call_fn(inputs, *args, **kwargs)
File "/usr/local/lib/python3.8/dist-packages/keras/utils/traceback_utils.
py", line 92, in error_handler
    return fn(*args, **kwargs)
File "/usr/local/lib/python3.8/dist-packages/keras/engine/sequential.py",
line 374, in call
    return super(Sequential, self).call(inputs, training=training, mask=mas
k)
File "/usr/local/lib/python3.8/dist-packages/keras/engine/functional.py",
line 458, in call
    return self._run_internal_graph(
File "/usr/local/lib/python3.8/dist-packages/keras/engine/functional.py",
line 596, in _run_internal_graph
    outputs = node.layer(*args, **kwargs)
File "/usr/local/lib/python3.8/dist-packages/keras/utils/traceback_utils.
py", line 64, in error_handler
    return fn(*args, **kwargs)
File "/usr/local/lib/python3.8/dist-packages/keras/engine/base_layer.py",
line 1014, in __call__
    outputs = call_fn(inputs, *args, **kwargs)
File "/usr/local/lib/python3.8/dist-packages/keras/utils/traceback_utils.
py", line 92, in error_handler
    return fn(*args, **kwargs)
File "/usr/local/lib/python3.8/dist-packages/keras/layers/core/dense.py",
line 235, in call
    outputs = self.activation(outputs)
File "/usr/local/lib/python3.8/dist-packages/keras/activations.py", line
311, in relu
    return backend.relu(x, alpha=alpha, max_value=max_value, threshold=thre
shold)
File "/usr/local/lib/python3.8/dist-packages/keras/backend.py", line 499
2, in relu
    x = tf.nn.relu(x)
Node: 'sequential/dense/Relu'

```

```
Matrix size-incompatible: In[0]: [10,196608], In[1]: [65536,128]
[[{{node sequential/dense/Relu}}]] [Op:__inference_train_function_15
34]
```

8. Accuracy

Plot accuracy value for train and validation

```
In [ ]: history.history.keys()
```

```
In [ ]: plt.plot(history.history['val_acc'])
plt.plot(history.history['acc'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()
```

```
In [ ]: score = model.evaluate(test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

9. RNN Architecture

The RNN uses the generic model , to show the accuracy of the data

```
In [24]: num_epochs = 2
model=tf.keras.models.Sequential()

model.add(layers.Input(shape=[256,256,3]))
model.add(layers.Rescaling(1/255))
model.add(layers.ConvLSTM1D(128,3,activation='relu'))
model.add(layers.BatchNormalization())
model.add(layers.Flatten())
model.add(layers.Dense(3,activation='softmax'))
```

In [25]: `model.summary()`

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
rescaling (Rescaling)	(None, 256, 256, 3)	0
conv_lstm1d (ConvLSTM1D)	(None, 254, 128)	201728
batch_normalization (Batch Normalization)	(None, 254, 128)	512
flatten_1 (Flatten)	(None, 32512)	0
dense_3 (Dense)	(None, 3)	97539
=====		
Total params: 299,779		
Trainable params: 299,523		
Non-trainable params: 256		

```
In [26]: model.compile(
    optimizer = 'adam',
    loss = 'categorical_crossentropy',
    metrics = ['accuracy']
)

model.fit(
    x = train,
    batch_size=1,
    epochs = 2,
    validation_data = (test),
    validation_steps = None,
    shuffle = False
)

pred_rnn = model.predict(test)
pred_rnn = np.argmax(pred_rnn, axis = 1)
print('\naccuracy: ', accuracy_score(test.labels, pred_rnn))
```


Epoch 1/2

7/1200 [.....] - ETA: 6:58:51 - loss: 1.1096 - accuracy: 0.2714

```

-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-26-a19101845148> in <module>
      5 )
      6
----> 7 model.fit(
      8     x = train,
      9     batch_size=1,

/usr/local/lib/python3.8/dist-packages/keras/utils/traceback_utils.py in error
r_handler(*args, **kwargs)
     62     filtered_tb = None
     63     try:
--> 64         return fn(*args, **kwargs)
     65     except Exception as e: # pylint: disable=broad-except
     66         filtered_tb = _process_traceback_frames(e.__traceback__)

/usr/local/lib/python3.8/dist-packages/keras/engine/training.py in fit(self,
x, y, batch_size, epochs, verbose, callbacks, validation_split, validation_da
ta, shuffle, class_weight, sample_weight, initial_epoch, steps_per_epoch, val
idation_steps, validation_batch_size, validation_freq, max_queue_size, worker
s, use_multiprocessing)
    1407         _r=1):
    1408             callbacks.on_train_batch_begin(step)
-> 1409             tmp_logs = self.train_function(iterator)
    1410             if data_handler.should_sync:
    1411                 context.async_wait()

/usr/local/lib/python3.8/dist-packages/tensorflow/python/util/traceback_util
s.py in error_handler(*args, **kwargs)
    148     filtered_tb = None
    149     try:
--> 150         return fn(*args, **kwargs)
    151     except Exception as e:
    152         filtered_tb = _process_traceback_frames(e.__traceback__)

/usr/local/lib/python3.8/dist-packages/tensorflow/python/eager/def_function.p
y in __call__(self, *args, **kws)
    913
    914     with OptionalXlaContext(self._jit_compile):
--> 915         result = self._call(*args, **kws)
    916
    917         new_tracing_count = self.experimental_get_tracing_count()

/usr/local/lib/python3.8/dist-packages/tensorflow/python/eager/def_function.p
y in _call(self, *args, **kws)
    945         # In this case we have created variables on the first call, so
we run the
    946         # defunned version which is guaranteed to never create variable
s.
--> 947         return self._stateless_fn(*args, **kws) # pylint: disable=not
-callable
    948     elif self._stateful_fn is not None:
    949         # Release the lock early so that multiple threads can perform t
he call

/usr/local/lib/python3.8/dist-packages/tensorflow/python/eager/function.py in

```

```

__call__(self, *args, **kwargs)
    2451         (graph_function,
    2452         filtered_flat_args) = self._maybe_define_function(args, kwargs)
    s)
-> 2453         return graph_function._call_flat(
    2454         filtered_flat_args, captured_inputs=graph_function.captured_inputs) # pylint: disable=protected-access
    2455

/usr/local/lib/python3.8/dist-packages/tensorflow/python/eager/function.py in
_call_flat(self, args, captured_inputs, cancellation_manager)
    1858         and executing_eagerly):
    1859         # No tape is watching; skip to running the function.
-> 1860         return self._build_call_outputs(self._inference_function.call(
    1861         ctx, args, cancellation_manager=cancellation_manager))
    1862         forward_backward = self._select_forward_and_backward_functions(

/usr/local/lib/python3.8/dist-packages/tensorflow/python/eager/function.py in
call(self, ctx, args, cancellation_manager)
    495         with _InterpolateFunctionError(self):
    496         if cancellation_manager is None:
--> 497         outputs = execute.execute(
    498         str(self.signature.name),
    499         num_outputs=self._num_outputs,

/usr/local/lib/python3.8/dist-packages/tensorflow/python/eager/execute.py in
quick_execute(op_name, num_outputs, inputs, attrs, ctx, name)
    52     try:
    53         ctx.ensure_initialized()
---> 54         tensors = pywrap_tfe.TFE_Py_Execute(ctx._handle, device_name, op_
name,
    55         inputs, attrs, num_outputs)
    56     except core._NotOkStatusException as e:

```

KeyboardInterrupt:

10. CNN Architecture

```
In [ ]: from keras.optimizers.optimizer_v1 import Optimizer

model=tf.keras.models.Sequential()

#First convolution layer
model.add(layers.Conv2D(16,(3,3),input_shape=[256,256,3],activation='relu'))
model.add(layers.MaxPooling2D(2,2))
#Second convolution layer
model.add(layers.Conv2D(32,(3,3),activation='relu'))
model.add(layers.MaxPooling2D(2,2))
#Third convolution layer
model.add(layers.Conv2D(64,(3,3),activation='relu'))
#Forth convolution layer
model.add(layers.Conv2D(64,(3,3),activation='relu'))
model.add(layers.MaxPooling2D(2,2))
#Forth convolution layer
model.add(layers.Conv2D(256,(3,3),activation='relu'))
model.add(layers.MaxPooling2D(2,2))
#Fifth convolution layer
model.add(layers.Conv2D(256,(3,3),activation='relu'))
model.add(layers.MaxPooling2D(2,2))
#Flatten the resiltis to feed into dense layer
model.add(layers.Flatten())
#Dropouts some neurons
model.add(layers.Dropout(0.5))
#64 neuron in fully connected layer
model.add(layers.Dense(64, activation='relu'))
#Three output neurons for three classes with softmax activation
model.add(layers.Dense(3, activation='softmax'))
```

```
In [ ]: # Summary
model.summary()
```

11. Evaluate on the test data

Compile and fit.

```
In [ ]: model.compile(optimizer='adam',loss='categorical_crossentropy',metrics=['accuracy'])
```

```
In [ ]: model.fit(train,
                  epochs=10,batch_size=10,
                  validation_data=(test),
                  callbacks=[callback],)
```

```
In [ ]: pred = model.predict(test)
pred_pre = np.argmax(pred, axis=1)
print('\naccuracy: ', accuracy_score(test.labels, pred_pre))

94/94 [=====] - 44s 465ms/step

accuracy:  0.32566666666666666
```

```
In [ ]: plt.plot(accuracy_score(test.labels,pred_pre),test)
```

12. Transfer learning

```
In [ ]: from tensorflow.keras.applications.efficientnet import EfficientNetB3
base_model = EfficientNetB3(include_top=False, input_shape=[256,256,3], weights="imagenet", pooling='max')
img_shape = [256,256,3]
num_epochs = 10

# init base model

fine_tune = 100
for layer in base_model.layers[:fine_tune]:
    layer.trainable = False

model_pre = tf.keras.models.Sequential()

model_pre.add(layers.Input(shape=img_shape))
model_pre.add(base_model)
model_pre.add(layers.BatchNormalization())
model_pre.add(layers.Dense(256, activation='relu'))
model_pre.add(layers.Dropout(rate=.2, seed=1234))
model_pre.add(layers.Dense(256, activation='relu'))
model_pre.add(layers.Dropout(rate=.2, seed=1234))
model_pre.add(layers.Dense(3, activation='softmax'))
```

```
In [ ]: model_pre.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
efficientnetb3 (Functional)	(None, 1536)	10783535
batch_normalization (Batch Normalization)	(None, 1536)	6144
dense_3 (Dense)	(None, 256)	393472
dropout_2 (Dropout)	(None, 256)	0
dense_4 (Dense)	(None, 256)	65792
dropout_3 (Dropout)	(None, 256)	0
dense_5 (Dense)	(None, 3)	771
Total params: 11,249,714		
Trainable params: 11,039,049		
Non-trainable params: 210,665		

```
In [ ]: model_pre.compile(
    optimizer = 'adam',
    loss = 'categorical_crossentropy',
    metrics = ['accuracy']
)

model_pre.fit(
    x = train,
    batch_size=10,
    epochs = 30,
    validation_data = (test),
    validation_steps = None,
    shuffle = False
)

pred_pre = model_pre.predict(test)
pred_pre = np.argmax(pred_pre, axis=1)
print('\naccuracy: ', accuracy_score(test.labels, pred_pre))
```

Epoch 1/30
375/375 [=====] - 289s 729ms/step - loss: 0.6167 - accuracy: 0.7593 - val_loss: 0.4757 - val_accuracy: 0.8463

Epoch 2/30
375/375 [=====] - 266s 710ms/step - loss: 0.3483 - accuracy: 0.8722 - val_loss: 0.1962 - val_accuracy: 0.9313

Epoch 3/30
375/375 [=====] - 265s 707ms/step - loss: 0.2356 - accuracy: 0.9157 - val_loss: 0.1300 - val_accuracy: 0.9513

Epoch 4/30
375/375 [=====] - 264s 705ms/step - loss: 0.2121 - accuracy: 0.9262 - val_loss: 0.1765 - val_accuracy: 0.9367

Epoch 5/30
375/375 [=====] - 264s 705ms/step - loss: 0.1831 - accuracy: 0.9356 - val_loss: 0.1360 - val_accuracy: 0.9527

Epoch 6/30
375/375 [=====] - 263s 701ms/step - loss: 0.1681 - accuracy: 0.9440 - val_loss: 0.1904 - val_accuracy: 0.9287

Epoch 7/30
375/375 [=====] - 263s 701ms/step - loss: 0.1503 - accuracy: 0.9492 - val_loss: 0.1169 - val_accuracy: 0.9550

Epoch 8/30
375/375 [=====] - 263s 702ms/step - loss: 0.1487 - accuracy: 0.9513 - val_loss: 0.1299 - val_accuracy: 0.9527

Epoch 9/30
375/375 [=====] - 263s 701ms/step - loss: 0.1557 - accuracy: 0.9477 - val_loss: 0.1314 - val_accuracy: 0.9590

Epoch 10/30
375/375 [=====] - 263s 700ms/step - loss: 0.1910 - accuracy: 0.9338 - val_loss: 2.7454 - val_accuracy: 0.6023

Epoch 11/30
375/375 [=====] - 263s 700ms/step - loss: 0.1753 - accuracy: 0.9420 - val_loss: 0.0912 - val_accuracy: 0.9667

Epoch 12/30
375/375 [=====] - 263s 701ms/step - loss: 0.1295 - accuracy: 0.9573 - val_loss: 0.1028 - val_accuracy: 0.9587

Epoch 13/30
375/375 [=====] - 263s 701ms/step - loss: 0.1292 - accuracy: 0.9574 - val_loss: 0.0789 - val_accuracy: 0.9733

Epoch 14/30
375/375 [=====] - 263s 700ms/step - loss: 0.1075 - accuracy: 0.9642 - val_loss: 0.0629 - val_accuracy: 0.9803

Epoch 15/30
375/375 [=====] - 263s 700ms/step - loss: 0.1246 - accuracy: 0.9602 - val_loss: 0.0782 - val_accuracy: 0.9697

Epoch 16/30
375/375 [=====] - 262s 699ms/step - loss: 0.1114 - accuracy: 0.9634 - val_loss: 0.1410 - val_accuracy: 0.9553

Epoch 17/30
375/375 [=====] - 262s 699ms/step - loss: 0.1175 - accuracy: 0.9616 - val_loss: 0.1017 - val_accuracy: 0.9663

Epoch 18/30
375/375 [=====] - 262s 698ms/step - loss: 0.1099 - accuracy: 0.9615 - val_loss: 0.1110 - val_accuracy: 0.9627

Epoch 19/30
375/375 [=====] - 261s 697ms/step - loss: 0.0989 - accuracy: 0.9678 - val_loss: 0.1201 - val_accuracy: 0.9623


```

Epoch 20/30
375/375 [=====] - 261s 697ms/step - loss: 0.0999 - a
ccuracy: 0.9667 - val_loss: 0.1223 - val_accuracy: 0.9603
Epoch 21/30
375/375 [=====] - 261s 695ms/step - loss: 0.0985 - a
ccuracy: 0.9677 - val_loss: 0.1191 - val_accuracy: 0.9610
Epoch 22/30
375/375 [=====] - 261s 696ms/step - loss: 0.1088 - a
ccuracy: 0.9651 - val_loss: 0.1318 - val_accuracy: 0.9543
Epoch 23/30
375/375 [=====] - 261s 696ms/step - loss: 0.0887 - a
ccuracy: 0.9702 - val_loss: 0.1228 - val_accuracy: 0.9583
Epoch 24/30
375/375 [=====] - 261s 697ms/step - loss: 0.1406 - a
ccuracy: 0.9532 - val_loss: 0.0705 - val_accuracy: 0.9750
Epoch 25/30
375/375 [=====] - 261s 697ms/step - loss: 0.0964 - a
ccuracy: 0.9697 - val_loss: 0.0725 - val_accuracy: 0.9797
Epoch 26/30
375/375 [=====] - 261s 696ms/step - loss: 0.0780 - a
ccuracy: 0.9746 - val_loss: 0.0606 - val_accuracy: 0.9763
Epoch 27/30
375/375 [=====] - 261s 695ms/step - loss: 0.0971 - a
ccuracy: 0.9678 - val_loss: 0.0814 - val_accuracy: 0.9773
Epoch 28/30
375/375 [=====] - 261s 696ms/step - loss: 0.0972 - a
ccuracy: 0.9692 - val_loss: 0.0993 - val_accuracy: 0.9723
Epoch 29/30
375/375 [=====] - 261s 696ms/step - loss: 0.0766 - a
ccuracy: 0.9752 - val_loss: 0.0802 - val_accuracy: 0.9770
Epoch 30/30
375/375 [=====] - 261s 694ms/step - loss: 0.0786 - a
ccuracy: 0.9743 - val_loss: 0.0871 - val_accuracy: 0.9760
94/94 [=====] - 49s 501ms/step

accuracy: 0.3363333333333333

```

```
In [ ]: base_model.summary()
```

13. Analysis

The data is to classify between three kinds of footwear: boots, shoes and sandle. The model CNN and RNN got the same tesable summary shows the tesability of the data is possible. CNN is giving slight better accuracy compare to RNN. Thoguh I tried to run RNN with increse epoch but the RAM was used up so could not execute or taking too long. The best should be the pre_model from the transfer leaning. It is giving the best accuracy at the moment with epoch of 30. Though it is still a low number it is high than CNN and RNN . The low batch_size did help but due to pre-runed the RAM & GPU was already low. The transfer learning tunes the data which is expected to give more accurate but for some reason it still low of .335 but greater than RNN and CNN. The model are functioning the issue is lack of RAM and GPU limitation.